

In [1]:

```
import os
from pathlib import Path

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt

import torch
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from torch import nn

from tqdm.auto import tqdm, trange

from transformers import AdamW

BATCH_SIZE = 64
```

In [2]:

```
!rm -rf PoS-Tagging
!git clone https://github.com/Janluke0/PoS-Tagging/
os.chdir('PoS-Tagging')
out_dir = Path('/kaggle/working/')
out_dir.mkdir(exist_ok=True)
```

Cloning into 'PoS-Tagging'...

remote: Enumerating objects: 86, done.

remote: Counting objects: 100% (86/86), done.

remote: Compressing objects: 100% (55/55), done.

remote: Total 86 (delta 30), reused 80 (delta 24), pack-reused 0

Unpacking objects: 100% (86/86), 15.25 MiB | 9.59 MiB/s, done.

In [3]:

```
from model.transformers.italian import ItBERTCasedPos, ItBERTUncasedPos
from model import train_model
from dataset import TWITADS
```

Common

this part should be added to the repo, it's time to go lighting

In [4]:

```
def train_model(model, dl_train, dl_test, cuda=False, lr=0.001, epochs=10, show_plots=False, save_dir=None):
    loss_function = nn.NLLLoss()
    optimizer = AdamW(model.parameters(), lr=lr, weight_decay=0.01)

    if cuda:
        model = model.cuda()
    if save_dir is not None:
        save_dir.mkdir(exist_ok=True)

    losses = []
    accuracies = []
    best_acc = 0
    best_loss = float('inf')
    pbar = trange(epochs)
    for epoch in pbar:
        model.train()
        for sample in tqdm(iter(dl_train), desc=f"Training {epoch}° epoch", leave=False)
```

```

x,m,y = sample['input_ids'],sample['attention_mask'], sample['labels']
if cuda:
    x, m, y = x.cuda(), m.cuda(), y.cuda()
optimizer.zero_grad()

tag_scores = model(input_ids=x,attention_mask=m)
loss = loss_function(tag_scores.transpose(1, 2),y)

loss.backward()
optimizer.step()

acc = []
los = []
## evaluation
model.eval()
with torch.no_grad():
    for sample in tqdm(iter(dl_test), desc=f"Eval {epoch}° epoch", leave=False):
        x,m,y = sample['input_ids'],sample['attention_mask'], sample['labels']
        if cuda:
            x, m, y = x.cuda(), m.cuda(), y.cuda()

        tag_scores = model(input_ids=x,attention_mask=m)
        if hasattr(tag_scores, 'logits'):
            tag_scores = tag_scores.logits

        loss = loss_function(tag_scores.transpose(1, 2),y)
        los.append(loss.cpu().item())

        acc.append(((tag_scores.argmax(2))==y)[m==1].float())

acc = torch.cat(acc).mean().item()
los = np.array(los).mean()

losses.append(los)
accuracies.append(acc)
#show epoch results
pbar.set_description(f"Loss:{los}\tAccuracy:{acc}")
if show_plots:
    plt.subplot(121)
    plt.title("Test loss")
    plt.plot(losses)

    plt.subplot(122)
    plt.title("Test accuracy")
    plt.plot(accuracies)
if save_dir is not None and acc >= best_acc:
    torch.save(model.state_dict(),save_dir/f"model_best_acc.pth")
if save_dir is not None and loss <= best_loss:
    torch.save(model.state_dict(),save_dir/f"model_best_loss.pth")

best_acc = max(acc,best_acc)
best_loss = max(los,best_loss)

return losses,accuracies

def show_pred(model, ds, i):
    REVTAG = {v:k for k,v in ds._TAGS.items()}
    model.cpu()
    sample = ds.collate([ds[i]])
    x,m,y = sample['input_ids'],sample['attention_mask'], sample['labels']
    with torch.no_grad():
        pred = model(input_ids=x,attention_mask=m)
        tkns = ds.tokenizer.convert_ids_to_tokens(x[0,1:-1])
        return list(zip(tkns,[REVTAG[v.item()] for v in pred[0].argmax(1)[1:-1]], [REVTAG[v.item()] for v in y[0][1:-1]]))

```

In [5]:

```

def collate_fn(batch):
    input_ids, token_type_ids, attention_mask, labels = [[] for _ in range(4)]
    for sample in batch:

```

```

input_ids.append(sample['input_ids'])
token_type_ids.append(sample['token_type_ids'])
attention_mask.append(sample['attention_mask'])
labels.append(sample['labels'])

d = {
    'input_ids':pad_sequence(input_ids,batch_first=True),
    'token_type_ids': pad_sequence(token_type_ids,batch_first=True),
    'labels':pad_sequence(labels, padding_value=-100,batch_first=True),
}
d['attention_mask'] = (d['labels']!=-100).float()*torch.ones(d['labels'].shape)
return d

```

In [6]:

```

def tokenize_and_align_labels(tokenizer, tokens, tags):
    tokens = list(tokens)
    tokenized_inputs = tokenizer(tokens, truncation=True, is_split_into_words=True)

    word_ids = tokenized_inputs.word_ids(batch_index=0) # Map tokens to their respective word.
    previous_word_idx = None
    label_ids = []
    for word_idx in word_ids: # Set the special tokens to -10
        if word_idx is None:
            label_ids.append(-100)
        elif word_idx != previous_word_idx: # Only label the first token of a given word.
            label_ids.append(tags[word_idx])
            previous_word_idx = word_idx
        else:
            label_ids.append(-100)

    tokenized_inputs["labels"] = label_ids

    return {k:torch.tensor(v) for k,v in tokenized_inputs.items()}

```

Cased model

In [7]:

```

cased_model = ItBERTCasedPos(23)

```

Some weights of the model checkpoint at dbmdz/bert-base-italian-cased were not used when initializing BertForTokenClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.bias', 'cls.seq_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.bias', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForTokenClassification were not initialized from the model checkpoint at dbmdz/bert-base-italian-cased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [8]:

```

tokenizer = ItBERTCasedPos.tokenizer()
ds_train = TWITADS('resampled_train',
                   lambda w:[w],
                   transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',

```

```
        lambda w:[w],
        transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
    )
```

In [9]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val   = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```

In [10]:

```
train_model
```


In [11]:

```
torch.manual_seed(42)
train_model(cased_model,dl_train,dl_val,cuda=torch.cuda.is_available(), lr=2e-5, epochs=
30, show_plots=True, save_dir=out_dir/"cased_model")
```

Out[11]:

```
([1.246960747241974,
 1.0487639367580415,
 1.0154193758964538,
 0.9969477176666259,
 0.9859574675559998,
 0.9781641244888306,
 0.9733147263526917,
 0.9744363844394683,
```

```
0.9719797492027282,  
0.9744706094264984,  
0.9706034243106842,  
0.9726258814334869,  
0.9762151360511779,  
0.9717868208885193,  
0.9796198666095733,  
0.9794528067111969,  
0.9824764490127563,  
0.9841284096240998,  
0.9887665688991547,  
0.9880777478218079,  
0.9851040422916413,  
0.9910544395446778,  
0.9932881772518158,  
0.9919820308685303,  
0.997039407491684,  
1.000197023153305,  
0.9995040416717529,  
0.9934559106826782,  
1.0007389307022094,  
1.0006577789783477],  
[0.8692137002944946,  
0.8857396841049194,  
0.8820080161094666,  
0.8707241415977478,  
0.8857396841049194,  
0.8899156451225281,  
0.8910706639289856,  
0.9011995196342468,  
0.9008440971374512,  
0.9054642915725708,  
0.9084851741790771,  
0.9134607315063477,  
0.9106175303459167,  
0.9190582036972046,  
0.922612190246582,  
0.9215459823608398,  
0.9233229756355286,  
0.9239449501037598,  
0.9258107542991638,  
0.9263438582420349,  
0.926432728767395,  
0.925011157989502,  
0.926876962184906,  
0.9266992807388306,  
0.9292759299278259,  
0.9285650849342346,  
0.9330075979232788,  
0.9323856234550476,  
0.934962272644043,  
0.930697500705719])
```



Uncased model

In [12]:

```
uncased_model = ItBERTUncasedPos(23)
```

Some weights of the model checkpoint at dbmdz/bert-base-italian-uncased were not used when initializing BertForTokenClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.bias', 'cls.seq_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.bias', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForTokenClassification were not initialized from the model checkpoint at dbmdz/bert-base-italian-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [13]:

```
tokenizer = ItBERTUncasedPos.tokenizer()
ds_train = TWITADS('resampled_train',
                  lambda w:[w],
                  transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',
                 lambda w:[w],
                 transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
```

In [14]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```

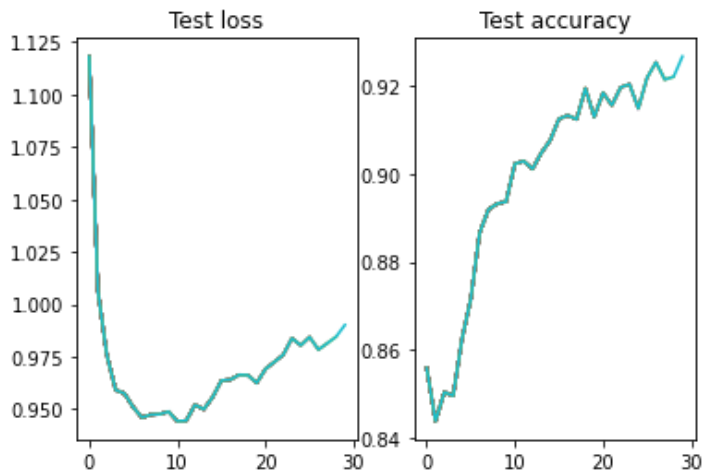
In [15]:

```
torch.manual_seed(42)
train_model(uncased_model,dl_train,dl_val,cuda=torch.cuda.is_available(), lr=2e-5, epochs=30, show_plots=True,save_dir=out_dir/"uncased_model")
```

Out[15]:

```
([1.1181639432907104,  
 1.006150245666504,  
 0.9759851455688476,  
 0.959184730052948,  
 0.9574857890605927,  
 0.9509188532829285,  
 0.9460362434387207,  
 0.947427386045456,  
 0.9476013720035553,  
 0.948854410648346,  
 0.9443119168281555,  
 0.9445123970508575,  
 0.9520338535308838,  
 0.9499574899673462,  
 0.9556913673877716,  
 0.963675731420517,  
 0.9640557587146759,  
 0.966399508714676,  
 0.9662407100200653,  
 0.9625177025794983,  
 0.9691667556762695,  
 0.9725211501121521,  
 0.9758700668811798,  
 0.9838366091251374,  
 0.9804231405258179,  
 0.9843933701515197,  
 0.978424996137619,  
 0.9814141809940338,  
 0.9845369637012482,  
 0.9901404619216919],  
 [0.8558478355407715,  
 0.8437610864639282,  
 0.8503376841545105,  
 0.8495378494262695,  
 0.8624244332313538,  
 0.8715783357620239,  
 0.8865979313850403,  
 0.8919302821159363,  
 0.8932633996009827,  
 0.8937077522277832,  
 0.9023284316062927,  
 0.9029505848884583,  
 0.9011731147766113,  
 0.9048168659210205,  
 0.9076608419418335,  
 0.912548840045929,  
 0.9133487343788147,  
 0.9124599695205688,  
 0.9194809794425964,  
 0.9129931926727295,  
 0.9185033440589905,  
 0.9156594276428223,  
 0.9107175010106760])
```

```
0.9197475910188788,
0.9204585552215576,
0.9149484038352966,
0.9219694137573242,
0.9254354238510132,
0.9215250611305237,
0.9221471548080444,
0.9267685413360596])
```



In [16]:

```
#remove repo from saved output
rm -rf /kagle/working/PoS-Tagging
```

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)