

In [1]:

```
import os
from pathlib import Path

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt

import torch
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from torch import nn

from tqdm.auto import tqdm, trange

from transformers import AdamW

BATCH_SIZE = 64
```

In [2]:

```
!rm -rf PoS-Tagging
!git clone https://github.com/Janluke0/PoS-Tagging/
os.chdir('PoS-Tagging')
out_dir = Path('/kaggle/working/')
out_dir.mkdir(exist_ok=True)
```

```
Cloning into 'PoS-Tagging'...
remote: Enumerating objects: 77, done.
remote: Counting objects: 100% (77/77), done.
remote: Compressing objects: 100% (49/49), done.
remote: Total 77 (delta 27), reused 71 (delta 21), pack-reused 0
Unpacking objects: 100% (77/77), 658.28 KiB | 837.00 KiB/s, done.
```

In [3]:

```
from model.transformers.italian import ItELECTRACasedPos, ItELECTRAXLCasePos
from dataset import TWITADS
```

## Common

this part should be added to the repo, it's time to go lighting

In [4]:

```
def train_model(model, dl_train, dl_test, cuda=False, lr=0.001, epochs=10, show_plots=False, save_dir=None):
    loss_function = nn.NLLLoss()
    optimizer = AdamW(model.parameters(), lr=lr, weight_decay=0.01)

    if cuda:
        model = model.cuda()
    if save_dir is not None:
        save_dir.mkdir(exist_ok=True)

    losses = []
    accuracies = []
    best_acc = 0
    best_loss = float('inf')
    pbar = trange(epochs)
    for epoch in pbar:
        model.train()
        for sample in tqdm(iter(dl_train), desc=f"Training {epoch}° epoch", leave=False):
            x, m, y = sample['input_ids'], sample['attention_mask'], sample['labels']
            if cuda:
                x, m, y = x.cuda(), m.cuda(), y.cuda()
            optimizer.zero_grad()

            tag_scores = model(input_ids=x, attention_mask=m)
            loss = loss_function(tag_scores.transpose(1, 2), y)

            loss.backward()
            optimizer.step()

        acc = []
        los = []
        ## evaluation
        model.eval()
        with torch.no_grad():
            for sample in tqdm(iter(dl_test), desc=f"Eval {epoch}° epoch", leave=False):
                x, m, y = sample['input_ids'], sample['attention_mask'], sample['labels']
                if cuda:
                    x, m, y = x.cuda(), m.cuda(), y.cuda()

                tag_scores = model(input_ids=x, attention_mask=m)
                if hasattr(tag_scores, 'logits'):
                    tag_scores = tag_scores.logits

                loss = loss_function(tag_scores.transpose(1, 2), y)
                los.append(loss.cpu().item())

                acc.append(((tag_scores.argmax(2)) == y)[m == 1].float())

        acc = torch.cat(acc).mean().item()
        los = np.array(los).mean()

    losses.append(los)
    accuracies.append(acc)
    #show epoch results
    pbar.set_description(f"Loss:{los}\tAccuracy:{acc}")
    if show_plots:
        plt.subplot(121)
```

```

plt.title("Test loss")
plt.plot(losses)

plt.subplot(122)
plt.title("Test accuracy")
plt.plot(accuracies)
if save_dir is not None and acc >= best_acc:
    torch.save(model.state_dict(), save_dir/f"model_best_acc.pth")
if save_dir is not None and loss <= best_loss:
    torch.save(model.state_dict(), save_dir/f"model_best_loss.pth")

best_acc = max(acc, best_acc)
best_loss = max(loss, best_loss)

return losses, accuracies

def show_pred(model, ds, i):
    REVTAG = {v:k for k,v in ds._TAGS.items()}
    model.cpu()
    sample = ds.collate([ds[i]])
    x,m,y = sample['input_ids'], sample['attention_mask'], sample['labels']
    with torch.no_grad():
        pred = model(input_ids=x, attention_mask=m)
        tkns = ds.tokenizer.convert_ids_to_tokens(x[0,1:-1])
    return list(zip(tkns, [REVTAG[v.item()] for v in pred[0].argmax(1)[1:-1]], [REVTAG[v.item()] for v in y[0][1:-1]]))

```

In [5]:

```

def collate_fn(batch):
    input_ids, token_type_ids, attention_mask, labels = [[] for _ in range(4)]
    for sample in batch:
        input_ids.append(sample['input_ids'])
        token_type_ids.append(sample['token_type_ids'])
        attention_mask.append(sample['attention_mask'])
        labels.append(sample['labels'])
    return {
        'input_ids': pad_sequence(input_ids, batch_first=True),
        'token_type_ids': pad_sequence(token_type_ids, batch_first=True),
        'labels': pad_sequence(labels, padding_value=-100, batch_first=True),
        'attention_mask': pad_sequence(attention_mask, batch_first=True),
    }

```

In [6]:

```
def tokenize_and_align_labels(tokenizer, tokens, tags):
    tokens = list(tokens)
    tokenized_inputs = tokenizer(tokens, truncation=True, is_split_into_words=True)

    word_ids = tokenized_inputs.word_ids(batch_index=0) # Map tokens to their respective word.
    previous_word_idx = None
    label_ids = []
    for word_idx in word_ids: # Set the special tokens to -1
        if word_idx is None:
            label_ids.append(-100)
        elif word_idx != previous_word_idx: # Only label the first token of a given word.
            label_ids.append(tags[word_idx])
            previous_word_idx = word_idx

    tokenized_inputs["labels"] = label_ids

    return {k:torch.tensor(v) for k,v in tokenized_inputs.items()}
```

## mc4 model

In [7]:

```
mc4_model = ItELECTRACasedPos(23)
```

Downloading:	686/686 [00:00<00:00,
100%	26.5kB/s]

Downloading:	419M/419M [00:18<00:00,
100%	24.8MB/s]

Some weights of the model checkpoint at dbmdz/electra-base-italian-mc4-cased-discriminator were not used when initializing ElectraForTokenClassification: ['discriminator\_predictions.dense.weight', 'discriminator\_predictions.dense\_prediction.weight', 'discriminator\_predictions.dense\_prediction.bias', 'discriminator\_predictions.dense.bias']

- This IS expected if you are initializing ElectraForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing ElectraForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of ElectraForTokenClassification were not initialized from the model checkpoint at dbmdz/electra-base-italian-mc4-cased-discriminator and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [8]:

```
tokenizer = ItELECTRACasedPos.tokenizer()
ds_train = TWITADS('resampled_train',
                   lambda w:[w],
                   transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',
                 lambda w:[w],
                 transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
```

Downloading:	59.0/59.0 [00:00<00:00,
100%	1.07kB/s]

Downloading:	230k/230k [00:00<00:00,
100%	320kB/s]

In [9]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```

In [10]:

```
torch.manual_seed(42)
train_model(mc4_model, dl_train, dl_val, cuda=torch.cuda.is_available(), lr=2e-5, epochs=4
0, show_plots=True, save_dir=out_dir/"mc4_model")
```

Loss:1.6618884563446046	40/40
Accuracy:0.9054461121559143: 100%	[23:13<00:00, 34.32s/it]
Training 0° epoch: 100%	91/91 [00:35<00:00, 3.16it/s]
Eval 0° epoch: 100%	10/10 [00:01<00:00, 7.38it/s]
Training 1° epoch: 100%	91/91 [00:32<00:00, 3.40it/s]
Eval 1° epoch: 100%	10/10 [00:01<00:00, 9.28it/s]
Training 2° epoch: 100%	91/91 [00:32<00:00, 3.18it/s]
Eval 2° epoch: 90%	9/10 [00:01<00:00, 8.38it/s]
Training 3° epoch: 100%	91/91 [00:32<00:00, 3.66it/s]
Eval 3° epoch: 90%	9/10 [00:01<00:00, 8.31it/s]
Training 4° epoch: 100%	91/91 [00:32<00:00, 3.42it/s]
Eval 4° epoch: 100%	10/10 [00:01<00:00, 8.72it/s]
Training 5° epoch: 100%	91/91 [00:32<00:00, 3.34it/s]
Eval 5° epoch: 90%	9/10 [00:01<00:00, 9.05it/s]
Training 6° epoch: 100%	91/91 [00:32<00:00, 3.47it/s]

Eval 6° epoch:	9/10 [00:01<00:00,
90%	8.87it/s]
Training 7° epoch:	91/91 [00:32<00:00,
100%	3.05it/s]
Eval 7° epoch:	10/10 [00:01<00:00,
100%	9.07it/s]
Training 8° epoch:	91/91 [00:32<00:00,
100%	3.61it/s]
Eval 8° epoch:	9/10 [00:01<00:00,
90%	8.67it/s]
Training 9° epoch:	91/91 [00:32<00:00,
100%	3.36it/s]
Eval 9° epoch:	9/10 [00:01<00:00,
90%	8.77it/s]
Training 10° epoch:	91/91 [00:32<00:00,
100%	3.15it/s]
Eval 10° epoch:	10/10 [00:01<00:00,
100%	8.94it/s]
Training 11° epoch:	91/91 [00:32<00:00,
100%	3.26it/s]
Eval 11° epoch:	10/10 [00:01<00:00,
100%	8.84it/s]
Training 12° epoch:	91/91 [00:31<00:00,
100%	3.25it/s]
Eval 12° epoch:	10/10 [00:01<00:00,
100%	8.53it/s]
Training 13° epoch:	91/91 [00:31<00:00,
100%	2.92it/s]



Eval 13° epoch: 100%	10/10 [00:01<00:00, 8.87it/s]
Training 14° epoch: 100%	91/91 [00:32<00:00, 3.12it/s]
Eval 14° epoch: 100%	10/10 [00:01<00:00, 8.84it/s]
Training 15° epoch: 100%	91/91 [00:31<00:00, 3.34it/s]
Eval 15° epoch: 100%	10/10 [00:01<00:00, 8.92it/s]
Training 16° epoch: 100%	91/91 [00:32<00:00, 3.46it/s]
Eval 16° epoch: 90%	9/10 [00:01<00:00, 8.86it/s]
Training 17° epoch: 100%	91/91 [00:32<00:00, 3.34it/s]
Eval 17° epoch: 90%	9/10 [00:01<00:00, 8.39it/s]
Training 18° epoch: 100%	91/91 [00:32<00:00, 3.30it/s]
Eval 18° epoch: 90%	9/10 [00:01<00:00, 8.68it/s]
Training 19° epoch: 100%	91/91 [00:32<00:00, 3.15it/s]
Eval 19° epoch: 100%	10/10 [00:01<00:00, 8.98it/s]
Training 20° epoch: 100%	91/91 [00:31<00:00, 3.20it/s]

Eval 20° epoch: 100%	10/10 [00:01<00:00, 8.81it/s]
Training 21° epoch: 100%	91/91 [00:32<00:00, 3.07it/s]
Eval 21° epoch: 100%	10/10 [00:01<00:00, 8.45it/s]
Training 22° epoch: 100%	91/91 [00:32<00:00, 3.17it/s]
Eval 22° epoch: 90%	9/10 [00:01<00:00, 8.57it/s]
Training 23° epoch: 100%	91/91 [00:31<00:00, 3.08it/s]
Eval 23° epoch: 100%	10/10 [00:01<00:00, 8.82it/s]
Training 24° epoch: 100%	91/91 [00:32<00:00, 3.25it/s]
Eval 24° epoch: 100%	10/10 [00:01<00:00, 9.15it/s]
Training 25° epoch: 100%	91/91 [00:32<00:00, 3.11it/s]
Eval 25° epoch: 100%	10/10 [00:01<00:00, 9.10it/s]
Training 26° epoch: 100%	91/91 [00:32<00:00, 3.42it/s]
Eval 26° epoch: 100%	10/10 [00:01<00:00, 9.30it/s]
Training 27° epoch: 100%	91/91 [00:32<00:00, 3.38it/s]


Eval 27° epoch: 100%	10/10 [00:01<00:00, 9.32it/s]
Training 28° epoch: 100%	91/91 [00:31<00:00, 3.54it/s]
Eval 28° epoch: 100%	10/10 [00:01<00:00, 8.79it/s]
Training 29° epoch: 100%	91/91 [00:32<00:00, 3.29it/s]
Eval 29° epoch: 90%	9/10 [00:01<00:00, 8.80it/s]
Training 30° epoch: 100%	91/91 [00:31<00:00, 3.30it/s]
Eval 30° epoch: 100%	10/10 [00:01<00:00, 9.14it/s]
Training 31° epoch: 100%	91/91 [00:31<00:00, 3.21it/s]
Eval 31° epoch: 90%	9/10 [00:01<00:00, 8.47it/s]
Training 32° epoch: 100%	91/91 [00:31<00:00, 3.50it/s]
Eval 32° epoch: 100%	10/10 [00:01<00:00, 8.88it/s]
Training 33° epoch: 100%	91/91 [00:32<00:00, 3.27it/s]
Eval 33° epoch: 90%	9/10 [00:01<00:00, 8.72it/s]
Training 34° epoch: 100%	91/91 [00:32<00:00, 3.31it/s]

Eval 34° epoch:	10/10 [00:01<00:00,
100%	8.98it/s]
Training 35° epoch:	91/91 [00:31<00:00,
100%	3.29it/s]
Eval 35° epoch:	10/10 [00:01<00:00,
100%	8.97it/s]
Training 36° epoch:	91/91 [00:32<00:00,
100%	3.27it/s]
Eval 36° epoch:	9/10 [00:01<00:00,
90%	8.60it/s]
Training 37° epoch:	91/91 [00:32<00:00,
100%	3.35it/s]
Eval 37° epoch:	10/10 [00:01<00:00,
100%	9.01it/s]
Training 38° epoch:	91/91 [00:31<00:00,
100%	3.14it/s]
Eval 38° epoch:	10/10 [00:01<00:00,
100%	9.30it/s]
Training 39° epoch:	91/91 [00:32<00:00,
100%	3.19it/s]
Eval 39° epoch:	10/10 [00:01<00:00,
100%	8.77it/s]

Out[10]:

```
([2.03640501499176,  
 1.7408401489257812,  
 1.6898200631141662,  
 1.664677894115448,  
 1.6550066471099854,  
 1.6488304495811463,  
 1.641623079776764,  
 1.6431013941764832,  
 1.643381667137146,  
 1.6387795448303222,  
 1.63668555021286,  
 1.6415698170661925,  
 1.6437358736991883,  
 1.6430535554885863,  
 1.6445921182632446,  
 1.6428797125816346,  
 1.6463431239128112,  
 1.6504170298576355,  
 1.6505876898765564,  
 1.650406789779663,  
 1.6526689291000367,  
 1.6517019391059875,  
 1.6518890380859375,  
 1.6510832071304322,  
 1.655223000049591,  
 1.6560324430465698,  
 1.6557697892189025,  
 1.6530117630958556,  
 1.6559265017509461,  
 1.6547646045684814,  
 1.657817280292511,  
 1.653559148311615,  
 1.653713059425354,  
 1.6519393086433412,  
 1.6572758078575134,  
 1.6560091733932496,  
 1.6590068221092225,  
 1.6578777313232422,  
 1.6527382850646972,  
 1.6618884563446046],  
 [0.8433812260627747,  
 0.8830790519714355,  
 0.8913741111755371,  
 0.8942872881889343,  
 0.8950278759002686,  
 0.900113582611084,  
 0.9036685824394226,  
 0.9037179946899414,  
 0.9038166999816895,  
 0.9058411121368408,  
 0.904952347278595,  
 0.9051004648208618,  
 0.9047054648399353,  
 0.9043104648590088,  
 0.9055448770523071,  
 0.906482994556427,  
 0.9058411121368408,  
 0.9051004648208618,  
 0.9069273471832275,
```

0.9059892296791077,  
0.9071248769760132,  
0.9067792296409607,  
0.905692994594574,  
0.9066311120986938,  
0.9078655242919922,  
0.9073717594146729,  
0.9055448770523071,  
0.9073717594146729,  
0.9067298769950867,  
0.9053967595100403,  
0.9058904647827148,  
0.907322347164154,  
0.9051498770713806,  
0.9058904647827148,  
0.907223641872406,  
0.9065817594528198,  
0.9042611122131348,  
0.907223641872406,  
0.9060386419296265,  
0.9054461121559143]])



**Uncased model**

In [11]:

```
xxl_model = ItELECTRAXLCasePos(23)
```

Downloading: 467/467 [00:00<00:00,  
100% 18.0kB/s]

Downloading: 419M/419M [00:18<00:00,  
100% 24.8MB/s]

Some weights of the model checkpoint at dbmdz/electra-base-italian-xxl-cased-discriminator were not used when initializing ElectraForTokenClassification: ['discriminator\_predictions.dense.weight', 'discriminator\_predictions.dense\_prediction.weight', 'discriminator\_predictions.dense\_prediction.bias', 'discriminator\_predictions.dense.bias']

- This IS expected if you are initializing ElectraForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing ElectraForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of ElectraForTokenClassification were not initialized from the model checkpoint at dbmdz/electra-base-italian-xxl-cased-discriminator and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [12]:

```
tokenizer = ItELECTRAXLCasePos.tokenizer()
ds_train = TWITADS('resampled_train',
                  lambda w:[w],
                  transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',
                lambda w:[w],
                transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
```

Downloading: 59.0/59.0 [00:00<00:00,  
100% 2.23kB/s]

Downloading: 230k/230k [00:00<00:00,  
100% 329kB/s]

In [13]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```

In [14]:

```
torch.manual_seed(42)
train_model(xx1_model,dl_train,dl_val,cuda=torch.cuda.is_available(), lr=2e-5, epochs=4
0, show_plots=True,save_dir=out_dir/"xx1_model")
```



Loss:1.6585968017578125	40/40
Accuracy:0.9071742296218872: 100%	[23:08<00:00, 34.28s/it]
Training 0° epoch: 100%	91/91 [00:34<00:00, 3.15it/s]
Eval 0° epoch: 100%	10/10 [00:01<00:00, 7.09it/s]
Training 1° epoch: 100%	91/91 [00:32<00:00, 3.36it/s]
Eval 1° epoch: 100%	10/10 [00:01<00:00, 9.28it/s]
Training 2° epoch: 100%	91/91 [00:32<00:00, 3.15it/s]
Eval 2° epoch: 90%	9/10 [00:01<00:00, 8.39it/s]
Training 3° epoch: 100%	91/91 [00:32<00:00, 3.64it/s]
Eval 3° epoch: 90%	9/10 [00:01<00:00, 8.30it/s]
Training 4° epoch: 100%	91/91 [00:32<00:00, 3.41it/s]
Eval 4° epoch: 100%	10/10 [00:01<00:00, 9.04it/s]
Training 5° epoch: 100%	91/91 [00:32<00:00, 3.37it/s]
Eval 5° epoch: 90%	9/10 [00:01<00:00, 9.01it/s]
Training 6° epoch: 100%	91/91 [00:32<00:00, 3.46it/s]

Eval 6° epoch:	9/10 [00:01<00:00,
90%	9.08it/s]
Training 7° epoch:	91/91 [00:32<00:00,
100%	3.03it/s]
Eval 7° epoch:	10/10 [00:01<00:00,
100%	9.03it/s]
Training 8° epoch:	91/91 [00:32<00:00,
100%	3.62it/s]
Eval 8° epoch:	9/10 [00:01<00:00,
90%	8.59it/s]
Training 9° epoch:	91/91 [00:32<00:00,
100%	3.35it/s]
Eval 9° epoch:	9/10 [00:01<00:00,
90%	8.75it/s]
Training 10° epoch:	91/91 [00:32<00:00,
100%	3.15it/s]
Eval 10° epoch:	10/10 [00:01<00:00,
100%	8.95it/s]
Training 11° epoch:	91/91 [00:32<00:00,
100%	3.28it/s]
Eval 11° epoch:	10/10 [00:01<00:00,
100%	8.96it/s]
Training 12° epoch:	91/91 [00:31<00:00,
100%	3.25it/s]
Eval 12° epoch:	10/10 [00:01<00:00,
100%	8.49it/s]
Training 13° epoch:	91/91 [00:31<00:00,
100%	2.92it/s]

Eval 13° epoch:	10/10 [00:01<00:00,
100%	8.86it/s]
Training 14° epoch:	91/91 [00:32<00:00,
100%	3.13it/s]
Eval 14° epoch:	10/10 [00:01<00:00,
100%	8.83it/s]
Training 15° epoch:	91/91 [00:31<00:00,
100%	3.33it/s]
Eval 15° epoch:	10/10 [00:01<00:00,
100%	8.95it/s]
Training 16° epoch:	91/91 [00:32<00:00,
100%	3.37it/s]
Eval 16° epoch:	9/10 [00:01<00:00,
90%	8.87it/s]
Training 17° epoch:	91/91 [00:32<00:00,
100%	3.32it/s]
Eval 17° epoch:	9/10 [00:01<00:00,
90%	8.40it/s]
Training 18° epoch:	91/91 [00:32<00:00,
100%	3.30it/s]
Eval 18° epoch:	9/10 [00:01<00:00,
90%	8.64it/s]
Training 19° epoch:	91/91 [00:32<00:00,
100%	3.16it/s]
Eval 19° epoch:	10/10 [00:01<00:00,
100%	8.96it/s]
Training 20° epoch:	91/91 [00:32<00:00,
100%	3.20it/s]

Eval 20° epoch:	10/10 [00:01<00:00,
100%	8.91it/s]
Training 21° epoch:	91/91 [00:32<00:00,
100%	3.07it/s]
Eval 21° epoch:	10/10 [00:01<00:00,
100%	8.65it/s]
Training 22° epoch:	91/91 [00:32<00:00,
100%	3.25it/s]
Eval 22° epoch:	9/10 [00:01<00:00,
90%	8.56it/s]
Training 23° epoch:	91/91 [00:32<00:00,
100%	3.13it/s]
Eval 23° epoch:	10/10 [00:01<00:00,
100%	8.78it/s]
Training 24° epoch:	91/91 [00:32<00:00,
100%	3.25it/s]
Eval 24° epoch:	10/10 [00:01<00:00,
100%	9.07it/s]
Training 25° epoch:	91/91 [00:32<00:00,
100%	3.11it/s]
Eval 25° epoch:	10/10 [00:01<00:00,
100%	9.07it/s]
Training 26° epoch:	91/91 [00:32<00:00,
100%	3.37it/s]
Eval 26° epoch:	10/10 [00:01<00:00,
100%	9.28it/s]
Training 27° epoch:	91/91 [00:32<00:00,
100%	3.36it/s]

Eval 27° epoch:	10/10 [00:01<00:00,
100%	9.26it/s]
Training 28° epoch:	91/91 [00:32<00:00,
100%	3.51it/s]
Eval 28° epoch:	10/10 [00:01<00:00,
100%	8.70it/s]
Training 29° epoch:	91/91 [00:32<00:00,
100%	3.28it/s]
Eval 29° epoch:	9/10 [00:01<00:00,
90%	8.84it/s]
Training 30° epoch:	91/91 [00:32<00:00,
100%	3.30it/s]
Eval 30° epoch:	10/10 [00:01<00:00,
100%	9.14it/s]
Training 31° epoch:	91/91 [00:31<00:00,
100%	3.22it/s]
Eval 31° epoch:	9/10 [00:01<00:00,
90%	8.62it/s]
Training 32° epoch:	91/91 [00:31<00:00,
100%	3.56it/s]
Eval 32° epoch:	10/10 [00:01<00:00,
100%	8.90it/s]
Training 33° epoch:	91/91 [00:32<00:00,
100%	3.28it/s]
Eval 33° epoch:	9/10 [00:01<00:00,
90%	8.67it/s]
Training 34° epoch:	91/91 [00:32<00:00,
100%	3.31it/s]

Eval 34° epoch:	10/10 [00:01<00:00,
100%	8.48it/s]
Training 35° epoch:	91/91 [00:32<00:00,
100%	3.12it/s]
Eval 35° epoch:	10/10 [00:01<00:00,
100%	8.82it/s]
Training 36° epoch:	91/91 [00:32<00:00,
100%	3.22it/s]
Eval 36° epoch:	9/10 [00:01<00:00,
90%	8.69it/s]
Training 37° epoch:	91/91 [00:32<00:00,
100%	3.34it/s]
Eval 37° epoch:	10/10 [00:01<00:00,
100%	9.01it/s]
Training 38° epoch:	91/91 [00:32<00:00,
100%	3.12it/s]
Eval 38° epoch:	10/10 [00:01<00:00,
100%	9.28it/s]
Training 39° epoch:	91/91 [00:32<00:00,
100%	3.19it/s]
Eval 39° epoch:	9/10 [00:01<00:00,
90%	8.53it/s]


Out[14]:

```
([1.8941421389579773,  
 1.697704815864563,  
 1.6630659103393555,  
 1.6501924037933349,  
 1.644134283065796,  
 1.6426794886589051,  
 1.6391411423683167,  
 1.638886296749115,  
 1.635600733757019,  
 1.6371947884559632,  
 1.632011342048645,  
 1.638020431995392,  
 1.64007648229599,  
 1.6401451706886292,  
 1.6420645475387574,  
 1.6433884739875793,  
 1.6441566348075867,  
 1.6422365188598633,  
 1.6461059331893921,  
 1.6471192479133605,  
 1.645556092262268,  
 1.648843777179718,  
 1.6501785159111022,  
 1.6515424489974975,  
 1.6503690958023072,  
 1.6521593928337097,  
 1.6550783038139343,  
 1.6489670991897583,  
 1.6582000613212586,  
 1.6511274218559264,  
 1.655085527896881,  
 1.6554555654525758,  
 1.6525050401687622,  
 1.6511088728904724,  
 1.655785608291626,  
 1.6527089476585388,  
 1.6493215084075927,  
 1.6608516097068786,  
 1.6534362316131592,  
 1.6585968017578125],  
 [0.8638226389884949,  
 0.8858440518379211,  
 0.8902384638786316,  
 0.8928554058074951,  
 0.8951266407966614,  
 0.8976941704750061,  
 0.900162935256958,  
 0.9028292298316956,  
 0.9036192297935486,  
 0.9029773473739624,  
 0.9046561121940613,  
 0.9048535823822021,  
 0.9020885825157166,  
 0.904211699962616,  
 0.9036685824394226,  
 0.9054461121559143,  
 0.9063842296600342,  
 0.9054461121559143,  
 0.906433641910553,
```

```

0.9051004648208618,
0.9082111120223999,
0.906482994556427,
0.9078655242919922,
0.905742347240448,
0.9075198769569397,
0.9069273471832275,
0.9052979946136475,
0.9074211120605469,
0.9043104648590088,
0.9059398770332336,
0.9058411121368408,
0.9067792296409607,
0.9053473472595215,
0.9047054648399353,
0.9058904647827148,
0.9076186418533325,
0.9060386419296265,
0.900113582611084,
0.9053473472595215,
0.9071742296218872])

```



In [15]:

```

#remove repo from saved output
rm -rf /kagle/working/PoS-Tagging

```

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS\_PARALLELISM=(true | false)