

In [1]:

```
import os
from pathlib import Path

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt

import torch
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from torch import nn

from tqdm.auto import tqdm, trange

from transformers import AdamW

BATCH_SIZE = 64
```

In [2]:

```
!rm -rf PoS-Tagging
!git clone https://github.com/Janluke0/PoS-Tagging/
os.chdir('PoS-Tagging')
out_dir = Path('/kaggle/working/')
out_dir.mkdir(exist_ok=True)
```

```
Cloning into 'PoS-Tagging'...
remote: Enumerating objects: 86, done.
remote: Counting objects: 100% (86/86), done.
remote: Compressing objects: 100% (55/55), done.
remote: Total 86 (delta 30), reused 80 (delta 24), pack-reused 0
Unpacking objects: 100% (86/86), 15.25 MiB | 8.80 MiB/s, done.
```

In [3]:

```
from model.transformers.italian import ItELECTRACasedPos, ItELECTRAXLCasePos
from dataset import TWITADS
```

Common

this part should be added to the repo, it's time to go lighting

In [4]:

```
def train_model(model, dl_train, dl_test, cuda=False, lr=0.001, epochs=10, show_plots=False, save_dir=None):
    loss_function = nn.NLLLoss()
    optimizer = AdamW(model.parameters(), lr=lr, weight_decay=0.01)

    if cuda:
        model = model.cuda()
    if save_dir is not None:
        save_dir.mkdir(exist_ok=True)

    losses = []
    accuracies = []
    best_acc = 0
    best_loss = float('inf')
    pbar = trange(epochs)
    for epoch in pbar:
        model.train()
        for sample in tqdm(iter(dl_train), desc=f"Training {epoch}° epoch", leave=False):
            x, m, y = sample['input_ids'], sample['attention_mask'], sample['labels']
```

```

        if cuda:
            x, m, y = x.cuda(), m.cuda(), y.cuda()
            optimizer.zero_grad()

            tag_scores = model(input_ids=x, attention_mask=m)
            loss = loss_function(tag_scores.transpose(1, 2), y)

            loss.backward()
            optimizer.step()

    acc = []
    los = []
    ## evaluation
    model.eval()
    with torch.no_grad():
        for sample in tqdm(iter(dl_test), desc=f"Eval {epoch}° epoch", leave=False):
            x, m, y = sample['input_ids'], sample['attention_mask'], sample['labels']
            if cuda:
                x, m, y = x.cuda(), m.cuda(), y.cuda()

            tag_scores = model(input_ids=x, attention_mask=m)
            if hasattr(tag_scores, 'logits'):
                tag_scores = tag_scores.logits

            loss = loss_function(tag_scores.transpose(1, 2), y)
            los.append(loss.cpu().item())

            acc.append(((tag_scores.argmax(2)) == y) [m==1].float())

    acc = torch.cat(acc).mean().item()
    los = np.array(los).mean()

    losses.append(los)
    accuracies.append(acc)
    #show epoch results
    pbar.set_description(f"Loss:{los}\tAccuracy:{acc}")
    if show_plots:
        plt.subplot(121)
        plt.title("Test loss")
        plt.plot(losses)

        plt.subplot(122)
        plt.title("Test accuracy")
        plt.plot(accuracies)
    if save_dir is not None and acc >= best_acc:
        torch.save(model.state_dict(), save_dir/f"model_best_acc.pth")
    if save_dir is not None and loss <= best_loss:
        torch.save(model.state_dict(), save_dir/f"model_best_loss.pth")

    best_acc = max(acc, best_acc)
    best_loss = max(los, best_loss)

    return losses, accuracies

def show_pred(model, ds, i):
    REVTAG = {v:k for k,v in ds._TAGS.items()}
    model.cpu()
    sample = ds.collate([ds[i]])
    x, m, y = sample['input_ids'], sample['attention_mask'], sample['labels']
    with torch.no_grad():
        pred = model(input_ids=x, attention_mask=m)
        tkns = ds.tokenizer.convert_ids_to_tokens(x[0,1:-1])
        return list(zip(tkns, [REVTAG[v.item()] for v in pred[0].argmax(1)[1:-1]], [REVTAG[v.item()] for v in y[0][1:-1]]))

```

In [5]:

```

def collate_fn(batch):
    input_ids, token_type_ids, attention_mask, labels = [[] for _ in range(4)]
    for sample in batch:
        input_ids.append(sample['input_ids'])

```

```

token_type_ids.append(sample['token_type_ids'])
attention_mask.append(sample['attention_mask'])
labels.append(sample['labels'])
d = {
    'input_ids':pad_sequence(input_ids,batch_first=True),
    'token_type_ids': pad_sequence(token_type_ids,batch_first=True),
    'labels':pad_sequence(labels, padding_value=-100,batch_first=True),
}
d['attention_mask'] = (d['labels']!=-100).float()*torch.ones(d['labels'].shape)
return d

```

In [6]:

```

def tokenize_and_align_labels(tokenizer, tokens, tags):
    tokens = list(tokens)
    tokenized_inputs = tokenizer(tokens, truncation=True, is_split_into_words=True)

    word_ids = tokenized_inputs.word_ids(batch_index=0) # Map tokens to their respective word.
    previous_word_idx = None
    label_ids = []
    for word_idx in word_ids: # Set the special tokens to -10
        if word_idx is None:
            label_ids.append(-100)
        elif word_idx != previous_word_idx: # Only label the first token of a given word.
            label_ids.append(tags[word_idx])
        previous_word_idx = word_idx

    tokenized_inputs["labels"] = label_ids

    return {k:torch.tensor(v) for k,v in tokenized_inputs.items()}

```

mc4 model

In [7]:

```
mc4_model = ItELECTRACasedPos(23)
```

Some weights of the model checkpoint at dbmdz/electra-base-italian-mc4-cased-discriminator were not used when initializing ElectraForTokenClassification: ['discriminator_predictions.dense.weight', 'discriminator_predictions.dense_prediction.weight', 'discriminator_predictions.dense.bias', 'discriminator_predictions.dense_prediction.bias']

- This IS expected if you are initializing ElectraForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing ElectraForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of ElectraForTokenClassification were not initialized from the model checkpoint at dbmdz/electra-base-italian-mc4-cased-discriminator and are newly initialized: ['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [8]:

```

tokenizer = ItELECTRACasedPos.tokenizer()
ds_train = TWITADS('resampled_train',
                  lambda w:[w],
                  transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',
                lambda w:[w],
                transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)

```

In [9]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```

In [10]:


```
torch.manual_seed(42)
train_model(mc4_model, dl_train, dl_val, cuda=torch.cuda.is_available(), lr=2e-5, epochs=40,
            show_plots=True, save_dir=out_dir/"mc4_model")
```

Out[10]:

```
([2.1494271993637084,
  1.7923119068145752,
  1.7144807696342468,
  1.6845481157302857,
  1.6607252101702141
```

1.668755491783141,
1.6575248479843139,
1.6495205521583558,
1.646745765209198,
1.6473207473754883,
1.643509840965271,
1.6396408677101135,
1.6404513835906982,
1.64548100233078,
1.6422146677970886,
1.6446046948432922,
1.6464330911636353,
1.6493518114089967,
1.6446463823318482,
1.645085048675537,
1.6472050666809082,
1.6457035899162293,
1.651600968837738,
1.648873221874237,
1.6514630794525147,
1.6510429978370667,
1.6511090755462647,
1.6538692712783813,
1.6521823167800904,
1.6569250464439391,
1.6521586775779724,
1.6570473194122315,
1.6530860304832458,
1.6495904922485352,
1.6495405912399292,
1.651939034461975,
1.6502442955970764,
1.653980541229248,
1.6541460633277894,
1.6584193110466003,
1.6597591519355774],
[0.847879946231842,
0.9255728125572205,
0.9404793381690979,
0.9472215175628662,
0.9493811130523682,
0.9544904232025146,
0.9585989117622375,
0.9616539478302002,
0.9592309594154358,
0.9608638286590576,
0.9608638286590576,
0.9624440670013428,
0.9603371024131775,
0.9624966979026794,
0.9633395075798035,
0.9637608528137207,
0.9647616744041443,
0.965657114982605,
0.9661838412284851,
0.9663418531417847,
0.965867817401886,
0.9664472341537476,
0.9662365317344666,
0.9658151268959045,
0.9671319723129272,
0.9652884006500244,
0.9670792818069458,
0.965025007724762,
0.9667105674743652,
0.9690281748771667,
0.9678167104721069,
0.9673953056335449,
0.9680801033973694,
0.9668159484863281,
0.9678694009780884,
0.9679747223854065,
0.9685511201272601

```
0.9683341391372881,
0.968238115310669,
0.9686068296432495,
0.967922031879425])
```



Uncased model

In [11]:

```
xxl_model = ItELECTRAXXLCasedPos(23)
```

Some weights of the model checkpoint at dbmdz/electra-base-italian-xxl-cased-discriminator were not used when initializing ElectraForTokenClassification: ['discriminator_predictions.dense.weight', 'discriminator_predictions.dense_prediction.weight', 'discriminator_predictions.dense.bias', 'discriminator_predictions.dense_prediction.bias']

- This IS expected if you are initializing ElectraForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing ElectraForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of ElectraForTokenClassification were not initialized from the model checkpoint at dbmdz/electra-base-italian-xxl-cased-discriminator and are newly initialized: ['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [12]:

```
tokenizer = ItELECTRAXXLCasedPos.tokenizer()
ds_train = TWITADS('resampled_train',
                  lambda w:[w],
                  transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
ds_val = TWITADS('resampled_validation',
                lambda w:[w],
                transform=lambda a,b: tokenize_and_align_labels(tokenizer,a,b)
)
```

In [13]:

```
dl_train = DataLoader(ds_train, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
dl_val = DataLoader(ds_val, shuffle=True, batch_size=BATCH_SIZE, collate_fn=collate_fn)
```


In [14]:


```
torch.manual_seed(42)
train_model(xxl_model,dl_train,dl_val,cuda=torch.cuda.is_available(), lr=2e-5, epochs=40,
            show_plots=True,save_dir=out_dir/"xxl_model")
```

Out[14]:

```
([1.944074845314026,  
  1.721768605709076,  
  1.6772701621055603,  
  1.6614771962165833,  
  1.6502413988113402,  
  1.647764301300049,  
  1.6428612351417542,  
  1.641972041130066,  
  1.6412869453430177,  
  1.6386998057365418,  
  1.636702871322632,  
  1.637635338306427,  
  1.641200637817383,  
  1.637195885181427,  
  1.6462345957756042,  
  1.6452369809150695,  
  1.6448238253593446.]
```

```
1.6476357221603393,  
1.6508209466934205,  
1.647178316116333,  
1.6420470118522643,  
1.6488866686820984,  
1.6487650513648986,  
1.6479997396469117,  
1.6493016242980958,  
1.6461843609809876,  
1.6514342427253723,  
1.647558343410492,  
1.6509463787078857,  
1.6502332091331482,  
1.6545352458953857,  
1.650391435623169,  
1.649366593360901,  
1.6525370240211488,  
1.6540325045585633,  
1.6521865010261536,  
1.6508914709091187,  
1.6557843804359436,  
1.6536250948905944,  
1.6531776309013366],  
[0.9078219532966614,  
0.9378983378410339,  
0.9429023265838623,  
0.9440611004829407,  
0.9476428627967834,  
0.9526468515396118,  
0.9579668045043945,  
0.960495114326477,  
0.9590203166007996,  
0.9591256380081177,  
0.9598103761672974,  
0.9606004953384399,  
0.9613379240036011,  
0.9625493884086609,  
0.9629181027412415,  
0.963286817073822,  
0.9622860550880432,  
0.9644456505775452,  
0.9631814956665039,  
0.9628127813339233,  
0.965025007724762,  
0.9657624363899231,  
0.9653410911560059,  
0.9629707932472229,  
0.9650776982307434,  
0.9670792818069458,  
0.9664472341537476,  
0.9653937220573425,  
0.9675006866455078,  
0.9664998650550842,  
0.9659731388092041,  
0.9664998650550842,  
0.9671319723129272,  
0.9671319723129272,  
0.9649196863174438,  
0.9667105674743652,  
0.9671319723129272,  
0.9648143649101257,  
0.9672372937202454,  
0.967922031879425])
```





In [15]:

```
#remove repo from saved output  
rm -rf /kagle/working/PoS-Tagging
```

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)